

Report

作业内容:

提交内容:

0、参考资料的学习笔记（阐述几个问题：什么是镜像？什么是容器？什么是NameSpace？什么是Cgroup？什么是Service？Service有什么用？Prometheus怎么工作的？）

1、自己编写的云应用代码github仓库

2、详细实验过程截图（证明自己的Service配置好了，证明自己的prometheus配置好了，能够拉到metrics了。）

3、README中简单阐述自己的代码逻辑

4、实验感想（可选）

上面的实验的评分逻辑:

0、问题回答是否正确且完整

1、代码注释是否清晰

2、metrics的设计是否合理且丰富

3、截图是否直观

4、README描述是否完整）

加分项（未完成加分项则满分90，加分项一个5分）:

0、截图证明配置了Grafana监控视图，并阐述如何配置

1、阐述Service API对象 对宿主机iptables的操作，截图查看宿主机iptables情况。

0.学习笔记

1.容器：docker容器就是用来运行镜像的小型虚拟机，或者说是一个轻量级的沙箱，Docker利用容器来运行和隔离应用。。

特性	虚拟机	容器
隔离级别	操作系统级	进程级
隔离策略	Hypervisor	CGroups
系统资源	5~15%	0~5%
启动时间	分钟级	秒级
镜像存储	GB-TB	KB-MB
集群规模	上百	上万
高可用策略	备份、容灾、迁移	弹性、负载、动态

2.镜像：类似于虚拟机镜像，镜像只是读的，是创建docker容器的基础。运行中的一个镜像就构成了一格容器。

3.NameSpace:Linux Namespace是Linux内核提供的一种资源隔离方案。处于不同 namespace 的进程拥有独立的全局系统资源，改变一个 namespace 中的系统资源只会影响当前 namespace 里的进程，对其他 namespace 中的进程没有影响。

名称	宏定义	隔离的资源
IPC	CLONE_NEWIPC	System V IPC(信号量、消息队列和共享内存) 和 POSIX message queues
Network	CLONE_NEWNET	Network devices, stacks, ports, etc(网络设备、网络栈、端口等).
Mount	CLONE_NEWNS	Mount points(文件系统挂载点)
PID	CLONE_NEWPID	Process IDs(进程编号)
User	CLONE_NEWUSER	User and group IDs(用户和用户组)
UTS	CLONE_NEWUTS	Hostname and NIS domain name(主机名与 NIS 域名)
Cgroup	CLONE_NEWCGROUP	Cgroup root directory(cgroup 的根目录)

4.Cgroups:Control Groups 是Linux内核的一个功能，用来限制、控制与分离一个进程组的资源。实现cgroups的主要目的是为不同用户层面的资源管理，提供一个统一化的接口。从单个进程的资源控制到操作系统层面的虚拟化。Cgroups提供了以下四大功能：

- 资源限制（Resource Limitation）：cgroups可以对进程组使用的资源总额进行限制。如设定应用运行时使用内存的上限，一旦超过这个配额就发出OOM（Out of Memory）。
- 优先级分配（Prioritization）：通过分配的CPU时间片数量及硬盘IO带宽大小，实际上就相当于控制了进程运行的优先级。
- 资源统计（Accounting）：cgroups可以统计系统的资源使用量，如CPU使用时长、内存用量等等，这个功能非常适用于计费。
- 进程控制（Control）：cgroups可以对进程组执行挂起、恢复等操作。

5.Services: Service是将运行在一组 [Pods](#) 上的应用程序公开为网络服务的抽象方法。Kubernetes `Service` 定义了这样一种抽象：逻辑上的一组 `Pod`，一种可以访问它们的策略——通常称为微服务。如果需要对应用程序进行访问，实际上是说需要对在某个node，某个pod上的应用程序进行访问，但由于在应用程序的生命周期内，pod可以是不同的，因此需要一个机制对其进行追踪，这就是service的用处。

6.Prometheus:Prometheus 项目工作的核心，是使用 Pull（抓取）的方式去搜集被监控对象的Metrics 数据（监控指标数据），然后，再把这些数据保存在一个 TSDB（时间序列数据库，比如 OpenTSDB、InfluxDB 等）当中，以便后续可以按照时间进行检索。Prometheus 剩下的组件就是用来配合这套机制的运行。Metrics 数据指3种数据：

- 宿主机的监控数据
- 来自于 Kubernetes 的 API Server、kubelet 等组件的 /metrics API
- Kubernetes 相关的监控数据

使用pull的方式由于类似一种轮询的机制，相比push方式实际上能降低阻塞的几率，提高可用性。

Prometheus提供了一个集成的规范化显示数据的机制，不必直接查看metrics相关文件。

reference

<https://zhuanlan.zhihu.com/p/53260098>

<https://www.kubernetes.org.cn/k8s>

<https://blog.csdn.net/ra681t58cjxsgckj31/article/details/104707642>

<https://en.wikipedia.org/wiki/Cgroups>

<https://docs.docker.com/engine/swarm/how-swarm-mode-works/services/>

<https://kubernetes.io/zh/docs/concepts/services-networking/service/>

<https://time.geekbang.org/column/article/72281>

2.实验过程截图

所有pod正常运行

```
request total 4
root@hourmor-virtual-machine:/home/hourmor/Downloads/example-ghm/src/metrics_version# curl 172.16.0.70:5565/abc
there is env Num. Computation succeeded
root@hourmor-virtual-machine:/home/hourmor/Downloads/example-ghm/src/metrics_version# kubectl get svc
NAME                TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
example-service     ClusterIP   10.99.7.139     <none>           80/TCP           16h
kubernetes           ClusterIP   10.96.0.1       <none>           443/TCP          16h
prometheus          NodePort    10.105.201.44   <none>           9090:31156/TCP   15h
root@hourmor-virtual-machine:/home/hourmor/Downloads/example-ghm/src/metrics_version# kubectl get pods -o wide
NAME                READY    STATUS    RESTARTS   AGE    IP              NODE                NOMINATED NODE    READINESS GATES
example-service-6c4d468bf6-pj85v    1/1      Running   0          111s   172.16.0.70     hourmor-virtual-machine    <none>             <none>
example-service-6c4d468bf6-txx6x    1/1      Running   0          110s   172.16.0.71     hourmor-virtual-machine    <none>             <none>
prometheus-7f57d48b8d-xg6v7         1/1      Running   2          12h    172.16.0.61     hourmor-virtual-machine    <none>             <none>
root@hourmor-virtual-machine:/home/hourmor/Downloads/example-ghm/src/metrics_version# curl 10.99.7.139/abc
there is env Num. Computation succeeded
root@hourmor-virtual-machine:/home/hourmor/Downloads/example-ghm/src/metrics_version# curl 10.99.7.139/abc
there is env Num. Computation succeeded
root@hourmor-virtual-machine:/home/hourmor/Downloads/example-ghm/src/metrics_version# curl 10.99.7.139/abc
there is env Num. Computation succeeded
root@hourmor-virtual-machine:/home/hourmor/Downloads/example-ghm/src/metrics_version# curl 10.99.7.139/metrics
# HELP bandwidth_rate rate of bandwidth lefted
# TYPE bandwidth_rate gauge
bandwidth_rate 0.4225
# HELP cpu_rate rate of cpu percent used.
# TYPE cpu_rate gauge
cpu_rate 3.5805626598454987
# HELP dis_rate rate of disk memory used.
# TYPE dis_rate gauge
dis_rate 37.749030655669976
# HELP go_gc_duration_seconds A summary of the pause duration of garbage collection cycles.
# TYPE go_gc_duration_seconds summary
go_gc_duration_seconds{quantile="0"} 0
go_gc_duration_seconds{quantile="0.25"} 0
go_gc_duration_seconds{quantile="0.5"} 0
go_gc_duration_seconds{quantile="0.75"} 0
go_gc_duration_seconds{quantile="1"} 0
go_gc_duration_seconds_sum 0
go_gc_duration_seconds_count 0
# HELP go_goroutines Number of goroutines that currently exist.
# TYPE go_goroutines gauge
go_goroutines 7
# HELP go_info Information about the Go environment.
# TYPE go_info gauge
```

Ln 91, Col 18 Tab Size: 4 UTF-8 LF Go

example-service 正常执行&&获取metrics

```
# HELP mem_rate rate of system memory used.
# TYPE mem_rate gauge
mem_rate 64.68306040966036
# HELP process_cpu_seconds_total Total user and system CPU time spent in seconds.
# TYPE process_cpu_seconds_total counter
process_cpu_seconds_total 0
# HELP process_max_fds Maximum number of open file descriptors.
# TYPE process_max_fds gauge
process_max_fds 1024
# HELP process_open_fds Number of open file descriptors.
# TYPE process_open_fds gauge
process_open_fds 85
# HELP process_resident_memory_bytes Resident memory size in bytes.
# TYPE process_resident_memory_bytes gauge
process_resident_memory_bytes 9.248768e+06
# HELP process_start_time_seconds Start time of the process since unix epoch in seconds.
# TYPE process_start_time_seconds gauge
process_start_time_seconds 1.59280999818e+09
# HELP process_virtual_memory_bytes Virtual memory size in bytes.
# TYPE process_virtual_memory_bytes gauge
process_virtual_memory_bytes 1.11130624e+09
# HELP process_virtual_memory_max_bytes Maximum amount of virtual memory available in bytes.
# TYPE process_virtual_memory_max_bytes gauge
process_virtual_memory_max_bytes -1
# HELP promhttp_metric_handler_requests_in_flight Current number of scrapes being served.
# TYPE promhttp_metric_handler_requests_in_flight gauge
promhttp_metric_handler_requests_in_flight 1
# HELP promhttp_metric_handler_requests_total Total number of scrapes by HTTP status code.
# TYPE promhttp_metric_handler_requests_total counter
promhttp_metric_handler_requests_total{code="200"} 0
promhttp_metric_handler_requests_total{code="500"} 0
promhttp_metric_handler_requests_total{code="503"} 0
# HELP request_latency_seconds Time spent in this service.
# TYPE request_latency_seconds histogram
request_latency_seconds_bucket{le="0.01"} 0
request_latency_seconds_bucket{le="0.02"} 0
request_latency_seconds_bucket{le="0.05"} 0
request_latency_seconds_bucket{le="0.1"} 0
request_latency_seconds_bucket{le="0.2"} 0
request_latency_seconds_bucket{le="0.5"} 0
request_latency_seconds_bucket{le="1"} 0
request_latency_seconds_bucket{le="2"} 3
request_latency_seconds_bucket{le="5"} 3
request_latency_seconds_bucket{le="10"} 3
request_latency_seconds_bucket{le="20"} 3
request_latency_seconds_bucket{le="30"} 3
request_latency_seconds_bucket{le="60"} 3
request_latency_seconds_bucket{le="120"} 3
request_latency_seconds_bucket{le="300"} 3
request_latency_seconds_bucket{le="+Inf"} 3
request_latency_seconds_sum 3.006253859
request_latency_seconds_count 3
# HELP request_total Number of request processed by this service.
# TYPE request_total counter
request_total 3
root@hourmor-virtual-machine:/home/hourmor/Downloads/example-ghm/src/deploy#
```

获取网络配置信息

```
Use 'apt autoremove' to remove them.
0 upgraded, 0 newly installed, 0 to remove and 147 not upgraded.
root@hourmor-virtual-machine:/home/hourmor/ex6/example-master/deploy# ifconfig
cni0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1450
    inet 172.16.0.1  netmask 255.255.255.0  broadcast 0.0.0.0
    inet6 fe80::ca7:38ff:fe10:a917  prefixlen 64  scopeid 0x20<link>
    ether 0e:a7:38:10:a9:17  txqueuelen 1000  (Ethernet)
    RX packets 11161  bytes 965486 (965.4 KB)
    RX errors 0  dropped 0  overruns 0  frame 0
    TX packets 9534  bytes 3656126 (3.6 MB)
    TX errors 0  dropped 0  overruns 0  carrier 0  collisions 0

docker0: flags=4099<UP,BROADCAST,MULTICAST>  mtu 1500
    inet 172.17.0.1  netmask 255.255.0.0  broadcast 172.17.255.255
    ether 02:42:3a:cc:91:5a  txqueuelen 0  (Ethernet)
    RX packets 0  bytes 0 (0.0 B)
    RX errors 0  dropped 0  overruns 0  frame 0
    TX packets 0  bytes 0 (0.0 B)
    TX errors 0  dropped 0  overruns 0  carrier 0  collisions 0

ens33: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
    inet 192.168.230.128  netmask 255.255.255.0  broadcast 192.168.230.255
    inet6 fe80::883e:094d:b14:8d69  prefixlen 64  scopeid 0x20<link>
    ether 00:0c:29:74:7e:a1  txqueuelen 1000  (Ethernet)
    RX packets 22463  bytes 31363839 (31.3 MB)
    RX errors 26  dropped 26  overruns 0  frame 0
    TX packets 6809  bytes 474530 (474.5 KB)
    TX errors 0  dropped 0  overruns 0  carrier 0  collisions 0
    device interrupt 19  base 0x2000

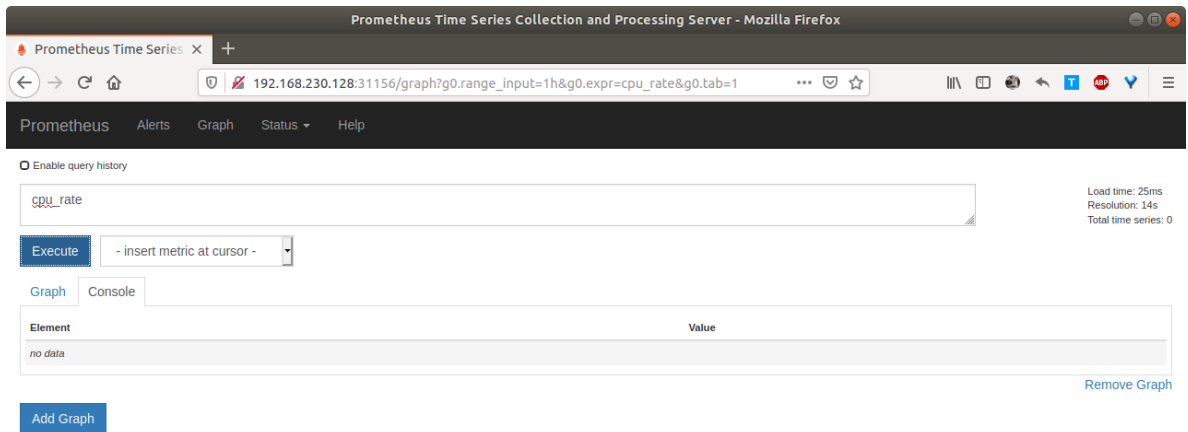
flannel.1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1450
    inet 172.16.0.0  netmask 255.255.255.255  broadcast 0.0.0.0
    inet6 fe80::40de:aaff:fe9e:ddeb  prefixlen 64  scopeid 0x20<link>
    ether 42:de:aa:9e:fd:eb  txqueuelen 0  (Ethernet)
    RX packets 0  bytes 0 (0.0 B)
    RX errors 0  dropped 0  overruns 0  frame 0
    TX packets 0  bytes 0 (0.0 B)
    TX errors 0  dropped 76  overruns 0  carrier 0  collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING>  mtu 65536
    inet 127.0.0.1  netmask 255.0.0.0
    inet6 ::1  prefixlen 128  scopeid 0x10<host>
    loop txqueuelen 1000  (Local Loopback)
    RX packets 160228  bytes 35208584 (35.2 MB)
    RX errors 0  dropped 0  overruns 0  frame 0
    TX packets 160228  bytes 35208584 (35.2 MB)
    TX errors 0  dropped 0  overruns 0  carrier 0  collisions 0

veth3f00c682: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1450
    inet6 fe80::6435:71ff:fea7:7dfc  prefixlen 64  scopeid 0x20<link>
    ether 66:35:71:a7:7d:fc  txqueuelen 0  (Ethernet)
    RX packets 47  bytes 22844 (22.8 KB)
    RX errors 0  dropped 0  overruns 0  frame 0
    TX packets 123  bytes 12257 (12.5 KB)
    TX errors 0  dropped 0  overruns 0  carrier 0  collisions 0
```

打开网页进行访问

这里发现实际上filter并没有筛选出关键字。打开发现是example-service在prometheus target对应状态为down。



mem_ra Highlight All Match Case Match Diacritics Whole Words Phrase not found

Prometheus Time Series Collection and Processing Server - Mozilla Firefox

Prometheus Time Series x + 192.168.230.128:31156/targets

Prometheus Alerts Graph Status Help

☐ Only unhealthy jobs

kubernetes-endpoints (1/8 up) [show less](#)

Endpoint	State	Labels	Last Scrape	Error
http://172.16.0.61:9090/metrics	UP	instance="172.16.0.61:9090"	14.958s ago	
http://172.16.0.63:53/metrics	DOWN	instance="172.16.0.63:53"	16.522s ago	context deadline exceeded
http://172.16.0.63:9153/metrics	DOWN	instance="172.16.0.63:9153"	17.031s ago	context deadline exceeded
http://172.16.0.65:53/metrics	DOWN	instance="172.16.0.65:53"	22.595s ago	context deadline exceeded
http://172.16.0.65:9153/metrics	DOWN	instance="172.16.0.65:9153"	20.544s ago	context deadline exceeded
http://172.16.0.70:5565/metrics	DOWN	instance="172.16.0.70:5565"	22.686s ago	context deadline exceeded
http://172.16.0.71:5565/metrics	DOWN	instance="172.16.0.71:5565"	10.208s ago	context deadline exceeded
http://192.168.230.128:6443/metrics	DOWN	instance="192.168.230.128:6443"	6.664s ago	server returned HTTP status 400 Bad Request

kubernetes-nodes (0/1 up) [show less](#)

Endpoint	State	Labels	Last Scrape	Error
http://192.168.230.128:10250/metrics	DOWN	instance="hourmor-virtual-machine"	11.684s ago	server returned HTTP status 400 Bad Request

kubernetes-pods (1/8 up) [show less](#)

Endpoint	State	Labels	Last Scrape	Error
http://172.16.0.61:9090/metrics	UP	instance="172.16.0.61:9090"	906ms ago	
http://172.16.0.63:53/metrics	DOWN	instance="172.16.0.63:53"	14.918s ago	context deadline exceeded
http://172.16.0.63:9153/metrics	DOWN	instance="172.16.0.63:9153"	21.173s ago	context deadline exceeded
http://172.16.0.65:53/metrics	DOWN	instance="172.16.0.65:53"	23.463s ago	context deadline exceeded
http://172.16.0.65:9153/metrics	DOWN	instance="172.16.0.65:9153"	16.798s ago	Get http://172.16.0.65:9153/metrics: dial tcp 172.16.0.65:9153: i/o timeout
http://172.16.0.70:5565/metrics	DOWN	instance="172.16.0.70:5565"	10.596s ago	context deadline exceeded
http://172.16.0.71:5565/metrics	DOWN	instance="172.16.0.71:5565"	18.321s ago	context deadline exceeded

mem_ra Highlight All Match Case Match Diacritics Whole Words Phrase not found

查找资料。说是config文件配置出错，原因可能有如下

```
scrape_interval: 5m
scrape_timeout: 1m

-----

tls_config:
  insecure_skip_verify: true
```

但我都试了，发现并没有如期生效，仍不知道是哪里出现了问题。心累，故放弃。

<https://stackoverflow.com/questions/49817558/context-deadline-exceeded-prometheus>

<https://github.com/prometheus/prometheus/issues/2459>

UPDATE

按照群里zzy同学得建议：

直接用 ./src/flannel.yaml 依次执行下列命令：

```
kubectl delete -f flannel.yaml
kubectl apply -f flannel.yaml
#####上面得命令更新得是kube-system kube-flannel
kubectl scale deployment prometheus --replicas=0
kubectl scale deployment prometheus --replicas=1
```

即可解决

Prometheus Time Series Collection and Processing Server - Mozilla Firefox

Prometheus Time Series

192.168.230.128:31156/graph?g0.range_input=1h&g0.expr=cpu_rate&g0.tab=1

DevDocs API Docume... WeChat for Web 题库 - 力扣 (LeetCode)... Welcome to the SICP ... CS 97SI: Introduction t... TensorBoard http://jupyterhub.vick... >>

Prometheus Alerts Graph Status Help

☐ Enable query history

cpu_rate

Execute - insert metric at cursor -

Load time: 19ms
Resolution: 14s
Total time series: 2

Graph Console

Element	Value
cpu_rate[instance="172.16.0.89:5565",job="kubernetes-endpoints"]	2.763819095478033
cpu_rate[instance="172.16.0.89:5565",job="kubernetes-pods"]	3.29949238578626

Remove Graph

Add Graph